

Class 7: Machine Learning 1

Alexander Liu (PID: 69026918)


Clustering

We will start with k-means clustering, one of the most prevalent of all clustering methods.

To get started let's make some data up:

```
hist(rnorm(10000, 3))
```

Histogram of rnorm(10000, 3)



```
tmp <- c(rnorm(30,3), rnorm(30, -3))  
x <- cbind(tmp, rev(tmp))  
x
```

tmp


[1,]	3.5216401	-3.6592951
[2,]	2.4133656	-3.3605628
[3,]	2.8358769	-4.0969981
[4,]	2.8609877	-3.9922001
[5,]	1.9699756	-2.3427514
[6,]	3.1121443	-2.7947423
[7,]	2.0715735	-2.4909884
[8,]	3.0217256	-0.7112393
[9,]	3.9218650	-2.0669755
[10,]	4.5889129	-2.3083050
[11,]	4.2635915	-3.3872075
[12,]	2.4253799	-2.6208772
[13,]	4.1752198	-3.5603367
[14,]	4.0816219	-3.0756043
[15,]	2.0386815	-1.7707004
[16,]	3.7402409	-2.2805179
[17,]	2.4043948	-3.7490486
[18,]	1.9497388	-2.5429718
[19,]	4.0450115	-2.0644694
[20,]	2.9891540	-4.0173051
[21,]	1.3627591	-3.3703716
[22,]	3.1047655	-4.0492415
[23,]	2.0295845	-2.8232469
[24,]	4.0025919	-3.2799778
[25,]	2.2786689	-2.9960491
[26,]	3.0375333	-4.7868841
[27,]	2.1179152	-2.9755287
[28,]	4.6209648	-1.7272890
[29,]	1.1153581	-3.0771897
[30,]	2.9077040	-3.5304005
[31,]	-3.5304005	2.9077040
[32,]	-3.0771897	1.1153581
[33,]	-1.7272890	4.6209648
[34,]	-2.9755287	2.1179152
[35,]	-4.7868841	3.0375333
[36,]	-2.9960491	2.2786689
[37,]	-3.2799778	4.0025919
[38,]	-2.8232469	2.0295845
[39,]	-4.0492415	3.1047655
[40,]	-3.3703716	1.3627591
[41,]	-4.0173051	2.9891540
[42,]	-2.0644694	4.0450115
[43,]	-2.5429718	1.9497388

```

[44,] -3.7490486  2.4043948
[45,] -2.2805179  3.7402409
[46,] -1.7707004  2.0386815
[47,] -3.0756043  4.0816219
[48,] -3.5603367  4.1752198
[49,] -2.6208772  2.4253799
[50,] -3.3872075  4.2635915
[51,] -2.3083050  4.5889129
[52,] -2.0669755  3.9218650
[53,] -0.7112393  3.0217256
[54,] -2.4909884  2.0715735
[55,] -2.7947423  3.1121443
[56,] -2.3427514  1.9699756
[57,] -3.9922001  2.8609877
[58,] -4.0969981  2.8358769
[59,] -3.3605628  2.4133656
[60,] -3.6592951  3.5216401

```

```
plot(x)
```



The main function in R for K-means clustering is called `kmeans()`.

```

k <- kmeans(x, centers=2, nstart=20)
k

```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

```
      tmp
1 -2.983643  2.966965
2  2.966965 -2.983643
```

Clustering vector:

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Within cluster sum of squares by cluster:

```
[1] 48.40206 48.40206
(between_SS / total_SS =  91.6 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"       "
```

Q1. How many points are in each cluster

```
k$size
```

```
[1] 30 30
```

Q2. The clustering result i.e. membership vector?

```
k$cluster
```

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```


Q3. Cluster centers

```
k$centers
```

```
      tmp
1 -2.983643  2.966965
2  2.966965 -2.983643
```


Q4. Make a plot of our data colored by clustering results with optionally the cluster centers shown.

```
plot(x, col=k$cluster, pch=16)
points(k$centers, col="blue", pch=15, cex=2)
```



Q5. Run kmeans again but cluster into 3 groups and plot the results like we did above.

```
l <- kmeans(x, centers=3, nstart=20)
plot(x, col=l$cluster, pch=16)
points(l$centers, col="blue", pch=15, cex=2)
```



K-means will always return a clustering result - even if there is no clear groupings.

#Hierarchical Clustering

main function: `hclust()`

```
hc <- hclust( dist(x) )  
hc
```


Call:

```
hclust(d = dist(x))
```

Cluster method : complete
Distance : euclidean
Number of objects: 60

```
plot(hc)  
abline(h=8, col="red")
```

Cluster Dendrogram




The function to get our clusters/groups from a `hclust` object is called `cutree()`

```
grps <- cutree(hc, h=8)  
grps
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2  
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

Q. Plot our hclust results in terms of our data colored by cluster membership.

```
plot(x, col=grps)
```



Principal Component Analysis (PCA)

Class 7 Lab

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)

dim(x)
```

```
[1] 17  5
```

```
## Preview the first 6 rows
head(x)
```

	X	England	Wales	Scotland	N.Ireland
1	Cheese	105	103	103	66
2	Carcass_meat	245	227	242	267
3	Other_meat	685	803	750	586

4	Fish	147	160	122	93
5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139

```
# Note how the minus indexing works
rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

```
dim(x)
```

```
[1] 17 4
```


```
x <- read.csv(url, row.names=1)
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?


I prefer the latter (row.names=1). If you accidentally run the former code, it could overwrite your processed dataset.


```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



Q3: Changing what optional argument in the above `barplot()` function results in the following plot?


```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)),)
```



Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

Lying on the diagonal means the value between two countries are similar.

```
pairs(x, col=rainbow(17), pch=16)
```



Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

There are more plots that are not on the diagonal, which indicates N.Ireland is more dissimilar to other countries.

PCA to the rescue


```
# Use the prcomp() PCA function
pca <- prcomp( t(x) )
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00


Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

```
# Plot PC1 vs PC2
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x))
```



Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

```
colors=c("orange", "red", "blue", "green")
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x), col=colors)
```



```
v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )
v
```


```
[1] 67 29 4 0
```

```
## or the second row here...
z <- summary(pca)
z$importance
```

	PC1	PC2	PC3	PC4
Standard deviation	324.15019	212.74780	73.87622	2.921348e-14
Proportion of Variance	0.67444	0.29052	0.03503	0.000000e+00
Cumulative Proportion	0.67444	0.96497	1.00000	1.000000e+00


Digging deeper (variable loadings)

```
## Lets focus on PC1 as it accounts for > 90% of variance
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,1], las=2 )
```



Q9: Generate a similar 'loadings plot' for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

```
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,2], las=2 )
```




Fresh_potatoes and Soft_drinks. They are the main drivers to “push” Wales and Scotland to negative or positive side, respectively.

Using ggplot for these figures


```
library(ggplot2)

df <- as.data.frame(pca$x)
df_lab <- tibble::rownames_to_column(df, "Country")

# Our first basic plot
ggplot(df_lab) +
  aes(PC1, PC2, col=Country) +
  geom_point()
```




```
ggplot(df_lab) +  
  aes(PC1, PC2, col=Country, label=Country) +  
  geom_hline(yintercept = 0, col="gray") +  
  geom_vline(xintercept = 0, col="gray") +  
  geom_point(show.legend = FALSE) +  
  geom_label(hjust=1, nudge_x = -10, show.legend = FALSE) +  
  expand_limits(x = c(-300,500)) +  
  xlab("PC1 (67.4%)") +  
  ylab("PC2 (28%)") +  
  theme_bw()
```




```
ld <- as.data.frame(pca$rotation)
ld_lab <- tibble::rownames_to_column(ld, "Food")


ggplot(ld_lab) +
  aes(PC1, Food) +
  geom_col()
```



```
ggplot(ld_lab) +
  aes(PC1, reorder(Food, PC1), bg=PC1) +
  geom_col() +
  xlab("PC1 Loadings/Contributions") +
  ylab("Food Group") +
  scale_fill_gradient2(low="purple", mid="gray", high="darkgreen", guide=NULL) +
  theme_bw()
```

The inbuilt biplot() can be useful for small datasets
`biplot(pca)`



2. PCA of RNA-seq data

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

```
      wt1 wt2 wt3 wt4 wt5 ko1 ko2 ko3 ko4 ko5
gene1 439 458 408 429 420 90  88  86  90  93
gene2 219 200 204 210 187 427 423 434 433 426
gene3 1006 989 1030 1017 973 252 237 238 226 210
gene4 783 792 829 856 760 849 856 835 885 894
gene5 181 249 204 244 225 277 305 272 270 279
gene6 460 502 491 491 493 612 594 577 618 638
```

Q10: How many genes and samples are in this data set?


```
nrow(rna.data)
```

```
[1] 100
```

100 genes

```
## Again we have to take the transpose of our data
pca <- prcomp(t(rna.data), scale=TRUE)

## Simple un polished plot of pc1 and pc2
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2")
```




```
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	9.6237	1.5198	1.05787	1.05203	0.88062	0.82545	0.80111
Proportion of Variance	0.9262	0.0231	0.01119	0.01107	0.00775	0.00681	0.00642
Cumulative Proportion	0.9262	0.9493	0.96045	0.97152	0.97928	0.98609	0.99251

	PC8	PC9	PC10
Standard deviation	0.62065	0.60342	3.345e-15
Proportion of Variance	0.00385	0.00364	0.000e+00
Cumulative Proportion	0.99636	1.00000	1.000e+00

```
plot(pca, main="Quick scree plot")
```



```
## Variance captured per PC  
pca.var <- pca$sdev2
```


```
## Percent variance is often more informative to look at  
pca.var.per <- round(pca.var/sum(pca.var)*100, 1)  
pca.var.per
```

```
[1] 92.6  2.3  1.1  1.1  0.8  0.7  0.6  0.4  0.4  0.0
```

```
barplot(pca.var.per, main="Scree Plot",  
        names.arg = paste0("PC", 1:10),
```

```
xlab="Principal Component", ylab="Percent Variation")
```


Scree Plot



```
## A vector of colors for wt and ko samples
colvec <- colnames(rna.data)
colvec[grep("wt", colvec)] <- "red"
colvec[grep("ko", colvec)] <- "blue"

plot(pca$x[,1], pca$x[,2], col=colvec, pch=16,
     xlab=paste0("PC1 (", pca.var.per[1], "%)"),
     ylab=paste0("PC2 (", pca.var.per[2], "%)"))


text(pca$x[,1], pca$x[,2], labels = colnames(rna.data), pos=c(rep(4,5), rep(2,5)))
```



```
library(ggplot2)


df <- as.data.frame(pca$x)

# Our first basic plot
ggplot(df) +
  aes(PC1, PC2) +
  geom_point()
```



```
# Add a 'wt' and 'ko' "condition" column
df$samples <- colnames(rna.data)
df$condition <- substr(colnames(rna.data),1,2)


p <- ggplot(df) +
  aes(PC1, PC2, label=samples, col=condition) +
  geom_label(show.legend = FALSE)
p
```



```
p + labs(title="PCA of RNASeq Data",
  subtitle = "PC1 clealy seperates wild-type from knock-out samples",
  x=paste0("PC1 (", pca.var.per[1], "%)"),
  y=paste0("PC2 (", pca.var.per[2], "%)"),
  caption="Class example data") +
  theme_bw()
```

PCA of RNASeq Data

PC1 clearly separates wild-type from knock-out samples



Class example data

```
loading_scores <- pca$rotation[,1]
```

```
## Find the top 10 measurements (genes) that contribute  
## most to PC1 in either direction (+ or -)
```

```
gene_scores <- abs(loading_scores)
```

```
gene_score_ranked <- sort(gene_scores, decreasing=TRUE)
```

```
## show the names of the top 10 genes
```

```
top_10_genes <- names(gene_score_ranked[1:10])
```

```
top_10_genes
```

```
[1] "gene100" "gene66"  "gene45"  "gene68"  "gene98"  "gene60"  "gene21"  
[8] "gene56"  "gene10"  "gene90"
```

```
sessionInfo()
```

R version 4.3.1 (2023-06-16)

Platform: aarch64-apple-darwin20 (64-bit)

Running under: macOS Ventura 13.5

Matrix products: default

BLAS: /Library/Frameworks/R.framework/Versions/4.3-arm64/Resources/lib/libRblas.0.dylib
LAPACK: /Library/Frameworks/R.framework/Versions/4.3-arm64/Resources/lib/libRlapack.dylib;

locale:

[1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8

time zone: America/Los_Angeles

tzcode source: internal

attached base packages:

[1] stats graphics grDevices utils datasets methods base

other attached packages:

[1] ggplot2_3.4.4

loaded via a namespace (and not attached):

[1] vctrs_0.6.4	cli_3.6.1	knitr_1.44	rlang_1.1.1
[5] xfun_0.40	generics_0.1.3	jsonlite_1.8.7	labeling_0.4.3
[9] glue_1.6.2	colorspace_2.1-0	htmltools_0.5.6.1	scales_1.2.1
[13] fansi_1.0.5	rmarkdown_2.25	grid_4.3.1	evaluate_0.22
[17] munsell_0.5.0	tibble_3.2.1	fastmap_1.1.1	yaml_2.3.7
[21] lifecycle_1.0.3	compiler_4.3.1	dplyr_1.1.3	pkgconfig_2.0.3
[25] farver_2.1.1	digest_0.6.33	R6_2.5.1	tidyselect_1.2.0
[29] utf8_1.2.4	pillar_1.9.0	magrittr_2.0.3	withr_2.5.1
[33] tools_4.3.1	gtable_0.3.4		